

Dynamic Indexes

Algorithmic Design

Alberto Casagrande

Email: `acasagrande@units.it`

a.y. 2020/2021

Motivations

A Simple Problem for Registry Office

Let us consider the registry office

For each newborn, they record a set of data e.g., name, birthday, parents, etc.

So, the registry data-set (hopefully) changes quite often

What if they frequently perform a birthday-based search on the data-set? E.g., Find all the baby born a given (variable) day?

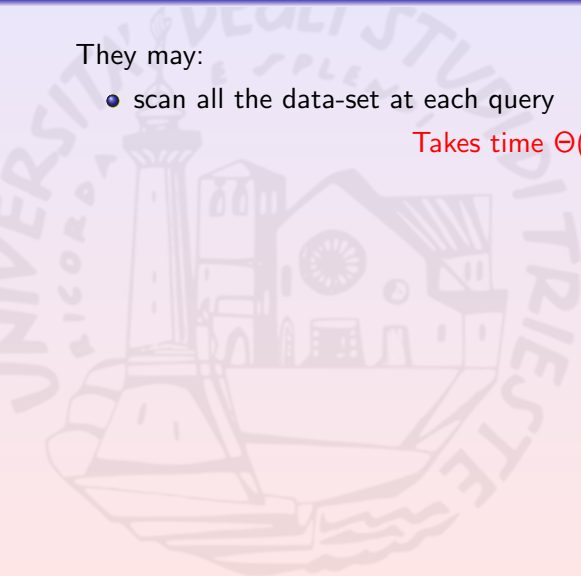
Some Possible Strategies

They may:



They may:

- Takes time $\Theta(n)$



They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $O(\log n)$ per query. What if they also want to perform searches by name?

They may:

- scan all the data-set at each query
Takes time $\Theta(n)$
- sort the data-set by birthday at each insertion
Takes time $\Theta(n)$ per insertion (Radix Sort) + $O(\log n)$ per query. What if they also want to perform searches by name?
- sort the data-set by birthday at each query
Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

They may:

- scan all the data-set at each query

Takes time $\Theta(n)$

- sort the data-set by birthday at each insertion

Takes time $\Theta(n)$ per insertion (Radix Sort) + $O(\log n)$ per query. What if they also want to perform searches by name?

- sort the data-set by birthday at each query

Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

- They may:
 - scan all the data-set at each query
Takes time $\Theta(n)$
 - sort the data-set by birthday at each insertion
Takes time $\Theta(n)$ per insertion (Radix Sort) + $O(\log n)$ per query. What if they also want to perform searches by name?
 - sort the data-set by birthday at each query
Takes time $\Theta(n + \log n) = \Theta(n)$ per query
(Radix Sort and DS)

The data-set often changes, thus, array is not the most suitable data structure to achieve this goal

The data-set often changes, thus, array is not the most suitable data structure to achieve this goal

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

A Dynamic Data Structure for Indexing

We need a data structure providing (efficient) support for:

- adding new data
- searching data
- removing data

We are aiming to build an **index** i.e., an auxiliary data structure to “efficiently” perform above operations

Binary Search Trees

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of x

x.key is the value stored in x i.e., its key

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of `x`

x.key is the value stored in x i.e., its key

x.left=NIL, then x misses the left child

Binary Search Tree Notation

We have already introduced trees as ADT.

As far as searching trees concern, the following notation holds

`x.left` and `x.right` are the left and right children of `x`, respectively

`x.parent` is the parent of x

x.key is the value stored in x i.e., its key

`x.left=NIL`, then `x` misses the left child

T.root is the root of the tree T and $T.root.parent = \text{NIL}$

```
def IS_ROOT(x):
    return x.parent==NIL
enddef

def IS_RIGHT_CHILD(x):
    return ¬IS_ROOT(x) and x.parent.right==x
enddef

def SIBLING(x):    # get x's sibling
    if IS_RIGHT_CHILD(x):
        return x.parent.left
    endif
    return x.parent.right
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd)

```
def CHILDHOOD_SIDE(x): # get x's side w.r.t.  
                        # its parent  
    if IS_RIGHT_CHILD(x):  
        return RIGHT  
    endif  
    return LEFT  
enddef  
  
def REVERSE_SIDE(side): # reverse the side  
    if side=LEFT:  
        return RIGHT  
    endif  
    return LEFT  
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd 2)

```
def GET_CHILD(x, side): # get x's child on side
    if side==LEFT:
        return x.left
    endif

    return x.right
enddef
```

Some Useful $\Theta(1)$ Functions (Cont'd 3)

```

def SET_CHILD(x, side, y):  # set x's child
    if side=LEFT:
        x.left ← y
    else:
        x.right ← y
    endif

    if y≠NIL:
        y.parent ← x
    endif
endef

```

Some Useful $\Theta(1)$ Functions (Cont'd 4)

```
def UNCLE(x):      #get x's uncle  
    return SIBLING(x.parent)  
enddef
```

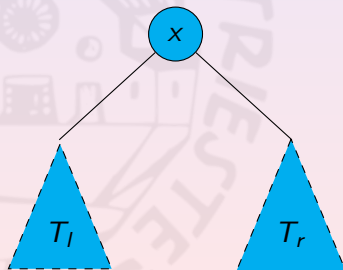
```
def GRANDPARENT(x):  # get x's granpa  
    return x.parent.parent  
enddef
```

```
def NEW_NODE(v):     # build a new node  
    x.key  $\leftarrow$  v  
    x.parent  $\leftarrow$  NIL  
    x.right  $\leftarrow$  NIL  
    x.left  $\leftarrow$  NIL  
enddef
```

Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:

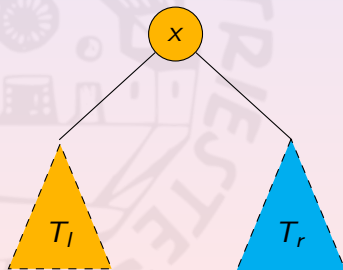
- all the keys belong to a totally ordered set w.r.t. \preceq



Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:

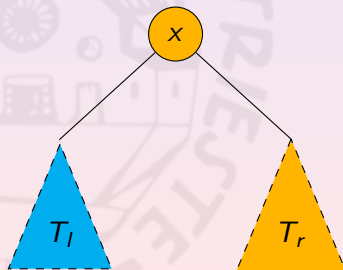
- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.\text{key} \preceq x.\text{key}$



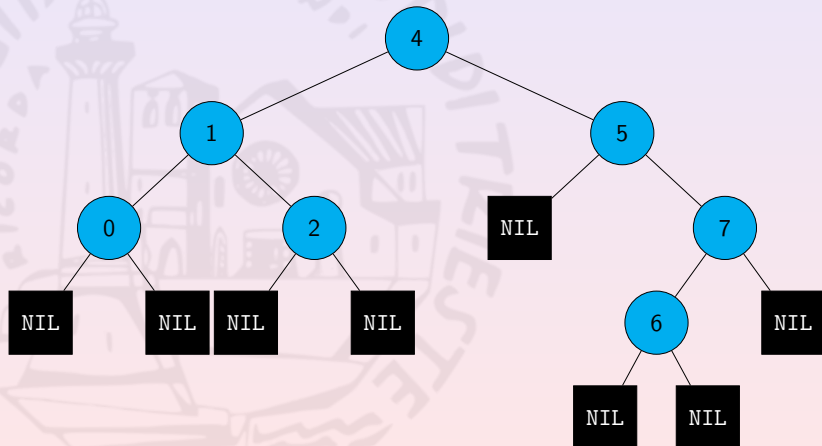
Binary Search Trees

A **Binary Search Tree (BST)** is a tree s.t.:


- all the keys belong to a totally ordered set w.r.t. \preceq
- if x_l is in the left sub-tree of x , then $x_l.key \preceq x.key$
- if x_r is in the right sub-tree of x , then $x.key \preceq x_r.key$



Binary Search Trees: an Example



Binary Search Trees: an Example



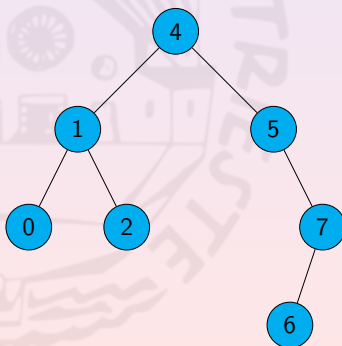
In-Order Walk

```
def INORDER_WALK_AUX(x):  
    if x ≠ NIL:  
        INORDER_WALK_AUX(x.left)  
        print x.key  
        INORDER_WALK_AUX(x.right)  
    endif  
endif  
  
def INORDER_WALK(T):  
    INORDER_WALK_AUX(T.root)  
endif
```

Searching for the Maximum/Minimum

Due to the BST property:

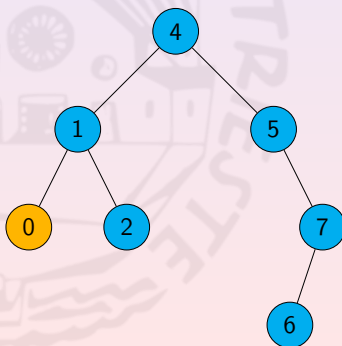
- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:

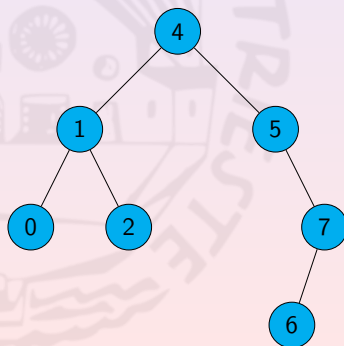
- the minimum key is contained by the first node on the leftmost branch that has not a left child



Searching for the Maximum/Minimum

Due to the BST property:

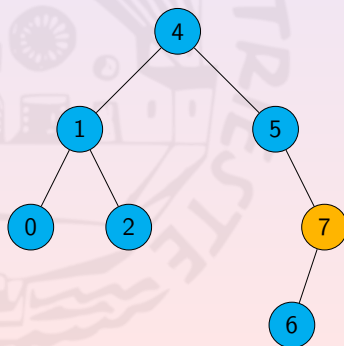
- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum

Due to the BST property:

- the minimum key is contained by the first node on the leftmost branch that has not a left child
- the maximum key is contained by the first node on the rightmost branch that has not a right child



Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_SUBTREE(x):  
    while x.left  $\neq$  NIL:  
        x  $\leftarrow$  x.left  
    endif  
    return x  
endif
```

```
def MAXIMUM_IN_SUBTREE(x):  
    while x.right  $\neq$  NIL:  
        x  $\leftarrow$  x.right  
    endif  
    return x  
endif
```

Searching for the Maximum/Minimum: Pseudo-Code

```
def MINIMUM_IN_SUBTREE(x):  
    while x.left ≠ NIL:  
        x ← x.left  
    endif  
    return x  
endif
```

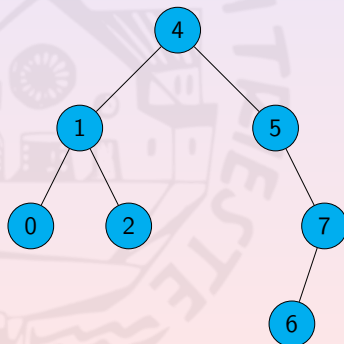
$O(h_T)$

```
def MAXIMUM_IN_SUBTREE(x):  
    while x.right ≠ NIL:  
        x ← x.right  
    endif  
    return x  
endif
```

$O(h_T)$

Successor of a Node

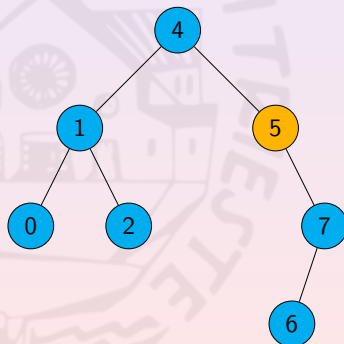
Due to the BST property, the successor w.r.t. \preceq of n is either



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

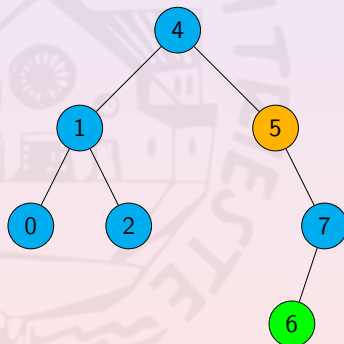
- the node containing the minimum in the right sub-tree of n



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

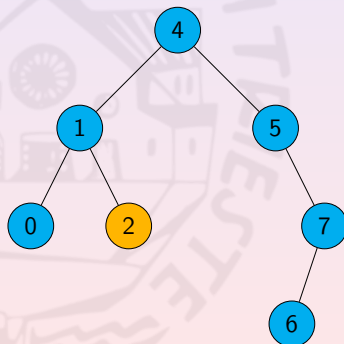
- the node containing the minimum in the right sub-tree of n or



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

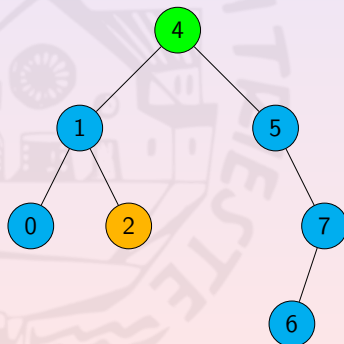
- the node containing the minimum in the right sub-tree of n or
- the nearest “right-ancestor” of n , if n has not right child



Successor of a Node

Due to the BST property, the successor w.r.t. \preceq of n is either

- the node containing the minimum in the right sub-tree of n or
- the nearest “right-ancestor” of n , if n has not right child



Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.right  $\neq$  NIL:  
        return MINIMUM_IN_SUBTREE(x.right)  
    endif  
  
    y  $\leftarrow$  x.parent  
    while y  $\neq$  NIL and IS_RIGHT_CHILD(x):  
        x  $\leftarrow$  y  
        y  $\leftarrow$  x.parent  
    endwhile  
  
    return y  
enddef
```


Successor of a Node: Pseudo-Code

```
def SUCCESSOR(x):  
    if x.right ≠ NIL:  
        return MINIMUM_IN_SUBTREE(x.right)  
    endif  
  
    y ← x.parent  
    while y ≠ NIL and IS_RIGHT_CHILD(x):  
        x ← y  
        y ← x.parent  
    endwhile  
  
    return y  
enddef
```

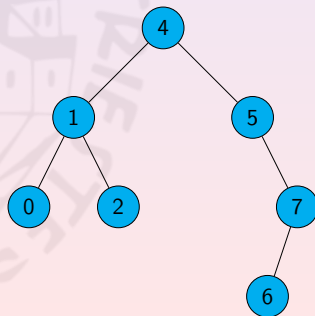
$O(h_T)$

Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- if $x.key \not\leq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- **if $x.key \not\leq v$, search on the left sub-tree**

Searching for 3

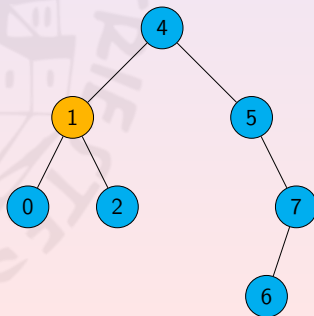


Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- **if $x.key \preceq v$, search on the right sub-tree**
- if $x.key \not\preceq v$, search on the left sub-tree

Searching for 3

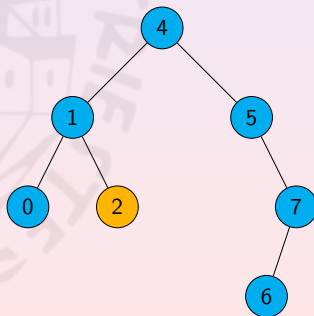


Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- **if $x.key \leq v$, search on the right sub-tree**
- if $x.key \not\leq v$, search on the left sub-tree

Searching for 3




Searching for a Value in a BST

To search for v , apply a dichotomic approach from the root x :

- if n is NIL or $x.key = v$, return n
- if $x.key \leq v$, search on the right sub-tree
- if $x.key \not\leq v$, search on the left sub-tree

Searching for 3



Searching for a Value in a BST: Pseudo-Code

```
def SEARCH_SUBTREE(x, v):  
    while x ≠ NIL:  
        if x.key ≤ v:  
            if v ≤ x.key:  
                return x  
            endif  
            x ← x.right  
        else:  
            x ← x.left  
        endif  
    endwhile  
enddef
```

Searching for a Value in a BST: Complexity

Each iteration performs $\Theta(1)$ operations

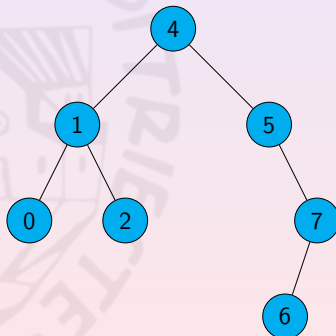
The # of iterations depends on the height h_T of T and on v

The algorithm takes time $O(h_T)$

Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

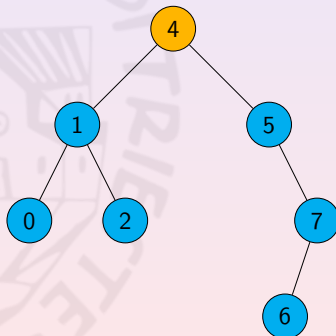
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

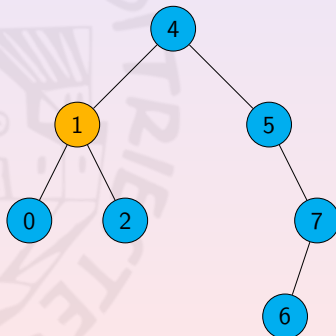
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

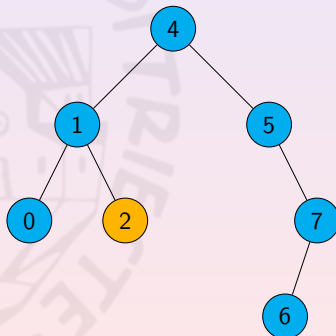
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key


Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

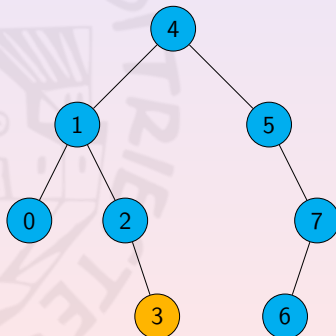
Inserting 3



Inserting a Value in a BST

Browse a branch of T and add a new leaf having v as key

Inserting 3



Inserting a Value in a BST: Pseudo-Code

```

def INSERT_BST(T,v): # v is the new value
    x ← T.root
    y ← NIL    # y is x's parent

    # search the right place for z
    while x ≠ NIL:
        y ← x
        if  $v \preceq x.key$ :
            if  $x.key \preceq v$ :
                return HANDLE.MULTI_INSERT(x,v)
            endif
            x ← x.left
        else:
            x ← x.right
        endif
    endwhile
  
```

Inserting a Value in a BST: Pseudo-Code

```

def INSERT_BST( $T, v$ ): # v is the new value
     $x \leftarrow T.\text{root}$ 
     $y \leftarrow \text{NIL}$    # y is x's parent

    # search the right place for z
    while  $x \neq \text{NIL}$ :
         $y \leftarrow x$ 
        if  $v \preceq x.\text{key}$ :
            if  $x.\text{key} \preceq v$ :
                return HANDLE.MULTIINSERT( $x, v$ )
            endif
             $x \leftarrow x.\text{left}$ 
        else:
             $x \leftarrow x.\text{right}$ 
        endif
    endwhile
  
```

$O(h_T)$

Inserting a Value in a BST: Pseudo-Code (Cont'd)

```

# attaching the new node
x ← NEW_NODE(v)
if T.root ≠ NIL:
    if v ≤ y.key:
        SET_CHILD(y, LEFT, x)
    else:
        SET_CHILD(y, RIGHT, x)
    endif
else:
    T.root ← x
endif

return x
enddef

```

Inserting a Value in a BST: Pseudo-Code (Cont'd)

```
# attaching the new node
x ← NEW_NODE(v)
if T.root ≠ NIL:
    if v ≤ y.key:
        SET_CHILD(y, LEFT, x)
    else:
        SET_CHILD(y, RIGHT, x)
    endif
else:
    T.root ← x
endif

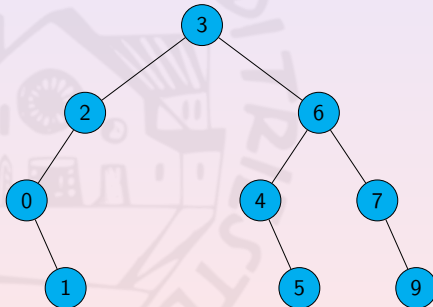
return x
enddef
```

 $\Theta(1)$

Removing a Key from a BST

Search the node x containing the key. Either

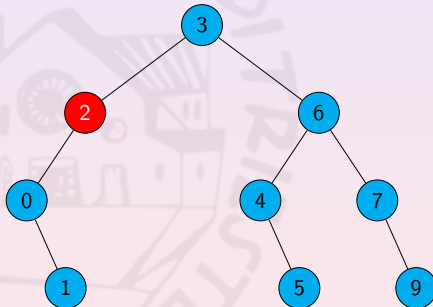
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

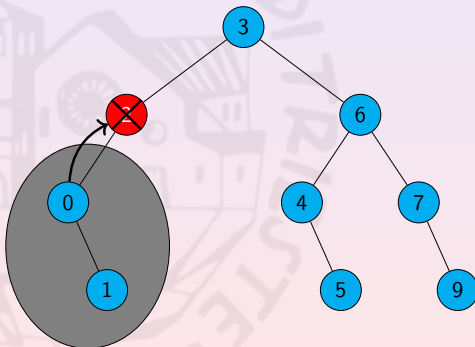
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

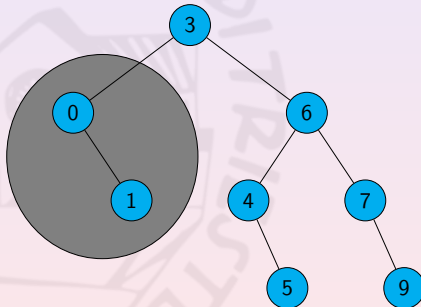
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

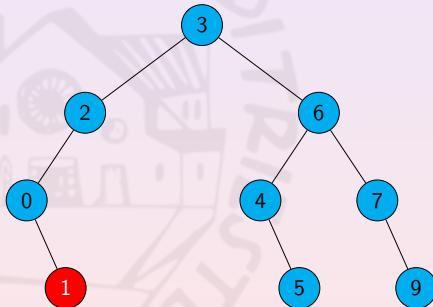
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

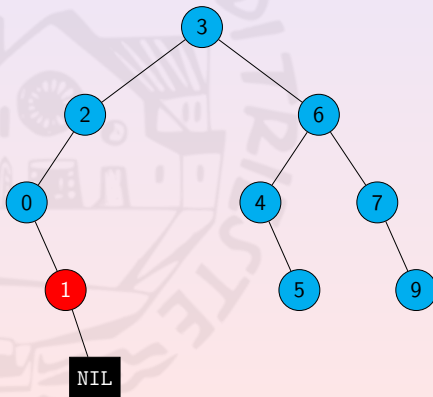
- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most



Removing a Key from a BST

Search the node x containing the key. Either

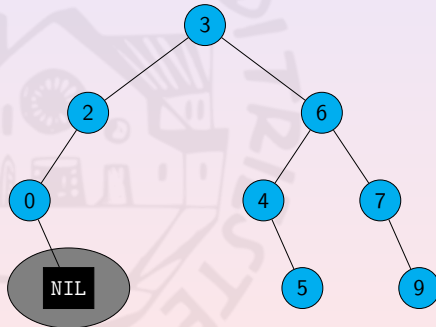
- x has one child at most or



Removing a Key from a BST

Search the node x containing the key. Either

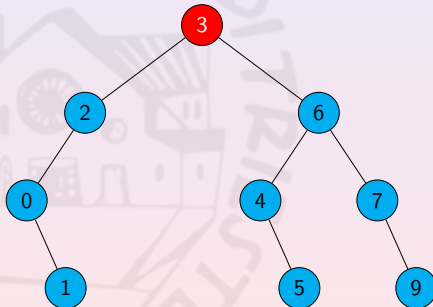
- x has one child at most or



Removing a Key from a BST

Search the node x containing the key. Either

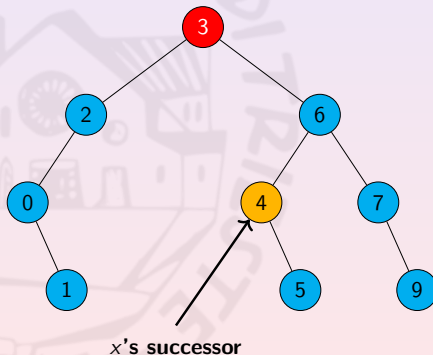
- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either


- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either

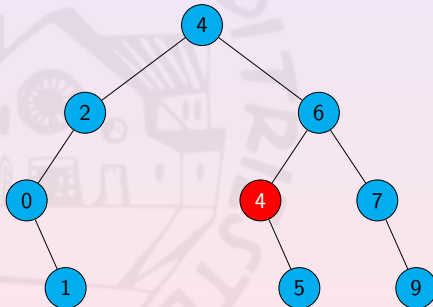
- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either

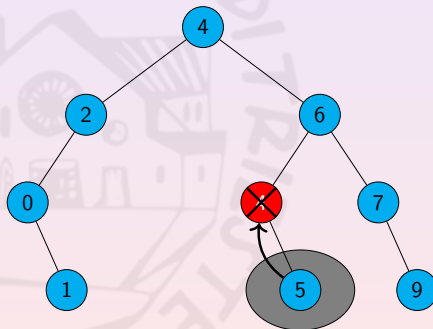
- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either

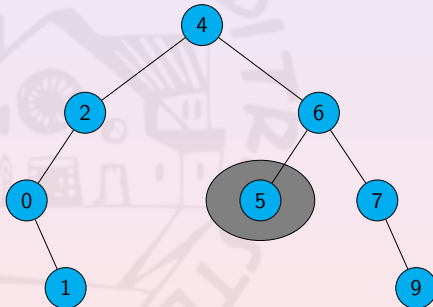
- x has one child at most or
- x has two children



Removing a Key from a BST

Search the node x containing the key. Either

- x has one child at most or
- x has two children



Removing a Key from a BST: Pseudo-Code

```

def TRANSPLANT(T,x,y): # replace x by y
    if IS_ROOT(x):
        T.root  $\leftarrow$  y

        if y  $\neq$  NIL:
            y.parent  $\leftarrow$  NIL
        endif
    else: # x has a parent
        x_side  $\leftarrow$  CHILDHOOD_SIDE(x)

        # attach y in place of x
        SET_CHILD(x.parent, x_side, y)
    endif
enddef

```

Removing a Key from a BST: Pseudo-Code (Cont'd)

```

def REMOVE_BST(T,x): # remove x.key from T and
                     # return a removed node
    if x.left=NIL:    # if x lacks of left child
        TRANSPLANT(T,x,x.right)
        return x
    endif

    if x.right=NIL:    # if x lacks of right child
        TRANSPLANT(T,x,x.left)
        return x
    endif

    y ← MINIMUM_IN_SUBTREE(x.right)
    x.key ← y.key
    return REMOVE_BST(T,y) # y lacks of left child
enddef

```

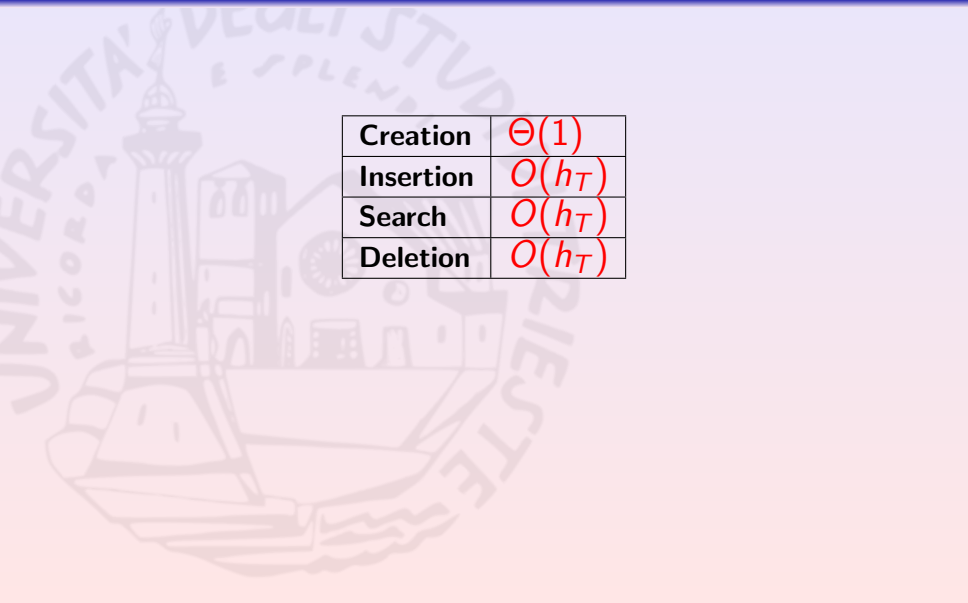
Remove a Node from a BST: Complexity

TRANSPLANT costs $\Theta(1)$, while MINIMUM_IN_TREE $O(h_T)$

So, if x has at most one child, removing it costs $\Theta(1)$

In the general case, removing x takes time $O(h_T)$

Summarizing BSTs...



Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

Summarizing BSTs...

Creation	$\Theta(1)$
Insertion	$O(h_T)$
Search	$O(h_T)$
Deletion	$O(h_T)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

Summarizing BSTs...

Creation	$\Theta(1)$
Insertion	$O(n)$
Search	$O(n)$
Deletion	$O(n)$

However, h_T may be equal to the number n of nodes e.g., keep inserting always the maximum

BSTs cost more than single-linked lists (insertion $\Theta(1)$)

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced?

Balancing BSTs

The minimum height for a binary tree having n nodes is $\lceil \log_2 n \rceil$

We aim to balance trees i.e., bring their heights to $O(\log n)$

How to do it?

GLOBALLY: balance the whole BST after each insertion/deletion
 $O(n)$

LOCALLY: balance only the “unbalanced” part of the tree
How to know if it is unbalanced? How to handle
branches' lengths?

The background of the slide features a large, faint watermark of the University of Trieste logo. The logo is circular and contains an illustration of a building with a dome and a tower, with the text "UNIVERSITA' DEGLI STUDI DI TRIESTE" and "FUNDATA 1868" around it.

Red-Black Trees

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node x , all the branches from x contain the same # of black nodes

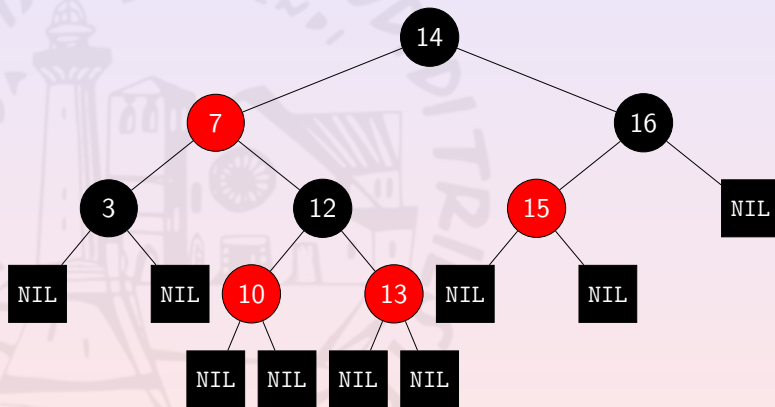
RBTs: Definition

Are BSTs satisfying the following conditions:

- each node is either a RED or a BLACK node
- the tree's root is BLACK
- all the leaves are BLACK NIL nodes
- all the RED nodes must have BLACK children
- for each node x , all the branches from x contain the same # of black nodes

$BH(x)$ will be the # of BLACK nodes below x in any branch

RBTs: An Example



Another Useful $\Theta(1)$ Function

```
def COLOR(x):  
    if x=NIL:  
        return BLACK  
    endif  
  
    return x.color  
enddef
```

How “Tall” Are RB-Trees?

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$

Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$

How “Tall” Are RB-Trees?

The ratio between x 's height and $BH(x)$ is topped by 2

Theorem (Heights of a RB-Tree)

Any RBT with n internal nodes has height at most $2 \log_2(n + 1)$


Proof Sketch:

- prove that the sub-tree rooted in x has at least $2^{BH(x)} - 1$ internal nodes
- $BH(x)$ is at least half of x 's height h then

$$n \geq 2^{h/2} - 1$$


Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



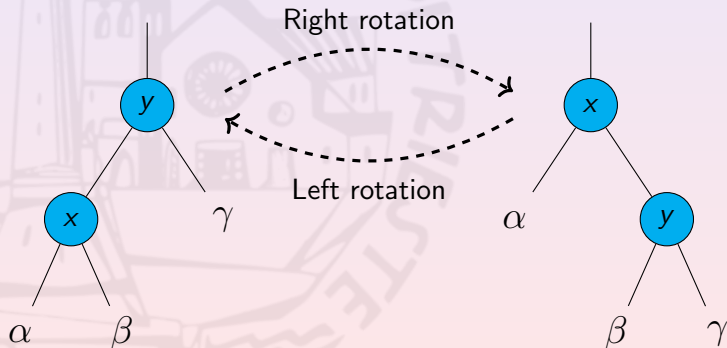
Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



Rotating a Sub-Tree

Rotations are operations on the tree structure. They preserve the BST property.



Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T,x,side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T,x,y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T,beta,x)

    SET_CHILD(x,r_side,beta) # move beta
enddef
  
```


Rotating a Sub-Tree: Pseudo-Code

```

def ROTATE(T, x, side):
    r_side ← REVERSE_SIDE(side)

    y ← GET_CHILD(x, r_side)
    TRANSPLANT(T, x, y)

    beta ← GET_CHILD(y, side)
    TRANSPLANT(T, beta, x)

    SET_CHILD(x, r_side, beta) # move beta
enddef

```

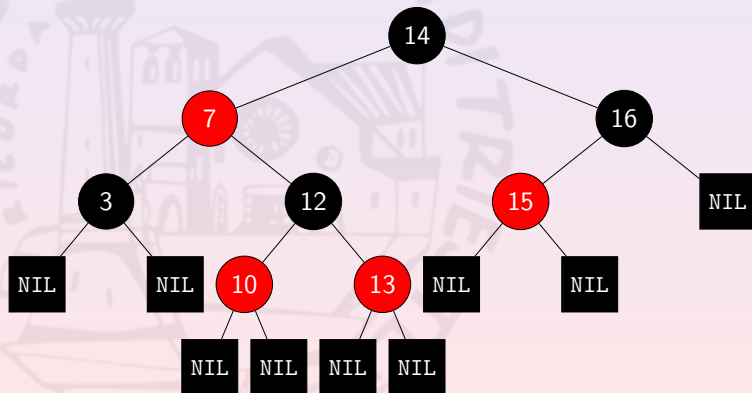
 $\Theta(1)$

Inserting a New Node

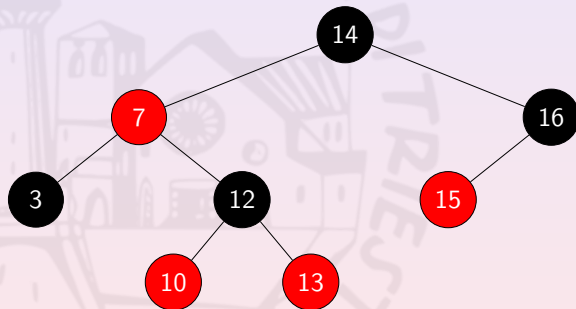
Requires:

- inserting as in BST
- RED-coloring the node
- fixing up RB-Tree properties

Inserting a New Node: Example



Inserting a New Node: Example




```

graph TD
    14((14)) --- 7((7))
    14 --- 16((16))
    7 --- 3((3))
    7 --- 12((12))
    12 --- 10((10))
    12 --- 13((13))
    10 --- 9((9))
    style 14 fill:#000,color:#fff
    style 7 fill:#f00,color:#fff
    style 16 fill:#000,color:#fff
    style 3 fill:#000,color:#fff
    style 12 fill:#000,color:#fff
    style 15 fill:#f00,color:#fff
    style 10 fill:#f00,color:#fff
    style 13 fill:#f00,color:#fff
    style 9 fill:#f00,color:#fff
    
```


Inserting a New Node: Example

x 's parent may be RED. How to fix it?




Inserting a New Node: Example

Case 1: x 's uncle (y) is RED...




Inserting a New Node: Example

Case 1: x 's uncle (y) is RED... RED-color x 's granpa and BLACK-color x 's parent and y .




Inserting a New Node: Example

Case 1: x 's uncle (y) is RED... RED-color x 's granpa and BLACK-color x 's parent and y . New x is former x 's granpa




Inserting a New Node: Example

Still facing problems, but x 's uncle is BLACK (not Case 1)




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side.




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side. Rotate on x 's parent on the opposite side.




Inserting a New Node: Example

Case 2: y is BLACK and y and x are on the same side. Rotate on x 's parent on the opposite side. New x is former x 's parent




Inserting a New Node: Example

Still facing problems, but




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1)




Inserting a New Node: Example

Still facing problems, but y is still BLACK (no Case 1) and x and y are on different sides (no Case 2)




Case 3: y is BLACK and y and x are on different sides.




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors



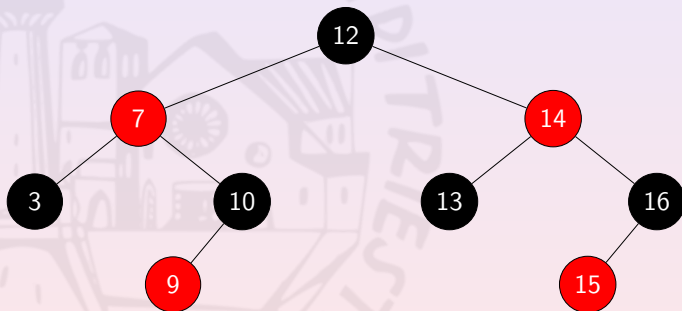
Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors and rotate on x 's granpa on y 's side




Inserting a New Node: Example

Case 3: y is BLACK and y and x are on different sides. Invert x 's pa and granpa colors and rotate on x 's granpa on y 's side



Inserting a New Node: Complexity

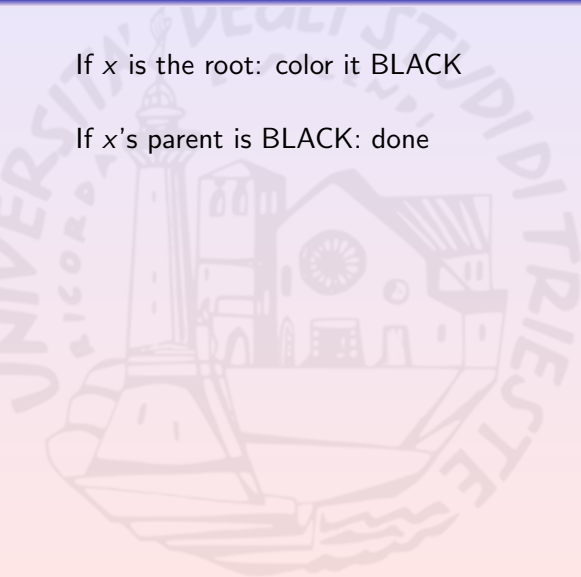
If x is the root: color it BLACK



Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done



Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Thus, we can always choose between:

Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Thus, we can always choose between:

Case 1 removes the problem or pushes it towards the root

Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Thus, we can always choose between:

Case 1 removes the problem or pushes it towards the root

Case 2 brings to Case 3

Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Thus, we can always choose between:

Case 1 removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

Inserting a New Node: Complexity

If x is the root: color it BLACK

If x 's parent is BLACK: done

If x 's parent is RED: it is not the root and x has an uncle

Thus, we can always choose between:

Case 1 removes the problem or pushes it towards the root

Case 2 brings to Case 3

Case 3 solves the problem

In the worst case, the algorithm keeps repeating Case 1 steps along the insertion branch and the complexity is $O(\log n)$

Inserting a New Node: Code

```
def INSERT_RBTREE(T, v):  
    x ← INSERT_BST(T, v)  
    x.color ← RED  
  
    FIX_INSERT_RBTREE(T, x)  
enddef  
  
def FIX_INSERT_RBT_CASE1(T, x):  
    UNCLE(x).color ← BLACK  
    x.parent.color ← BLACK  
    GRANDPARENT(x).color ← RED  
  
    return GRANDPARENT(x)  
endif
```

Inserting a New Node: Code (Cont'd)

```
def FIX_INSERT_RBT_CASE2(T,x):  
    x_side ← CHILDHOOD_SIDE(x)  
  
    p ← x.parent  
    ROTATE(T,p, REVERSE_SIDE(x_side))  
  
    return p  
endif
```

Inserting a New Node: Code (Cont'd 2)

```
def FIX_INSERT_RBT_CASE3(T, x):  
    g ← GRANDPARENT(z)  
  
    x.parent.color ← BLACK  
    g.color ← RED  
  
    x_side ← CHILDHOOD_SIDE(x)  
  
    ROTATE(T, g, REVERSE_SIDE(x_side))  
endif
```


Inserting a New Node: Code (Cont'd 3)

```

def FIX_INSERT_RBTREE(T,x):
    while (¬IS_ROOT(x) and
           (¬IS_ROOT(x.parent) or x.parent.color=RED)):
        if COLOR(UNCLE(x))=RED:
            z ← FIX_INSERT_RBT_CASE1(T,x)
        else:
            if (CHILDHOOD_SIDE(x) ≠
                CHILDHOOD_SIDE(x.parent)):
                z ← FIX_INSERT_RBT_CASE2(T,x)
            endif
            FIX_INSERT_RBT_CASE3(T,x)
        endif
    endwhile

    T.root.color ← BLACK
enddef

```

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x



Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved

If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

Removing a Key

Removing a key as in BST removes also a node y which is replaced by its former child x

If y was RED, the RB-Tree properties are preserved


If y was BLACK, the branches through x lost 1 BLACK node

In particular, $BH(x) = BH(w) - 1$ where w is x 's sibling

The fixing procedure iteratively balances BH on the sub-tree rooted on x 's parent

Removing a Key: Case 0

x is RED




Removing a Key: Case 0

x is RED


- BLACK-color x

$BH(x)$ is increased by 1 and the tree has been fixed




Removing a Key: Case 1

x 's sibling is RED



Removing a Key: Case 2

x 's sibling and nephews are BLACK





Removing a Key: Case 2

x 's sibling and nephews are BLACK


- RED-color x 's sibling

$BH(x)$ does not change, while the BLACK height of both x 's parent and sibling are decreased by 1



Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED




Removing a Key: Case 3

Among x 's sibling and nephews, only the nephew on x 's side is RED


- rotate x 's sibling on the opposite side w.r.t. x
- invert colors in both old and new siblings of x

The BLACK height of both x and x 's parent does not change



Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED




Removing a Key: Case 4

The x 's nephew on the opposite side w.r.t. x is RED


- switch colors of x 's parent and sibling
- BLACK-color the x 's nephew on the opposite side w.r.t. x
- rotate x 's parent on x 's side

The BLACK height of both x and x 's parent does not change



Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$



Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Removing a Key: Some Considerations

Each of the case transformation procedures takes time $\Theta(1)$

Both Case 0 and Case 4 transformation procedures fix the RB-Tree

Case 1 cannot occur twice in-row

Case 2 pushes the problem one step towards the root

If Case 2 occurs after Case 1, Case 0 occurs next

Case 3 transformation procedure brings to Case 4

Removing a Key: Some Considerations

In the worst case scenario, Case 2 is repeated until the problem is pushed up to the root ($O(\log n)$ times)

If this is the case, we have decreased the BLACK height of the tree and the problem is no more a problem

Removing a Key: Code

```
def REMOVE_RBTREE(T, z):  
    y ← REMOVE_BST(T, z)  
  
    if y.color = BLACK: # fix from its replacement  
        if y.left = NIL  
            x ← y.right  
        else:  
            x ← y.left  
        endif  
        FIX_REMOVE_RBTREE(T, x)  
    endif  
  
    return y  
enddef
```

Removing a Key: Code (Cont'd 2)

```
def FIX_REMOVE_RBT_CASE1(T, x):  
    SIBLING(x).color ← BLACK  
    x.parent.color ← RED  
  
    ROTATE(T, x, CHILDHOOD_SIDE(x))  
endif  
  
def FIX_REMOVE_RBT_CASE2(T, x):  
    SIBLING(x).color ← RED  
    return x.parent  
endif
```


Removing a Key: Code (Cont'd 3)

```
def FIX_REMOVE_RBT_CASE3(T,x):  
    x_side ← CHILDHOOD_SIDE(x)  
    r_side ← REVERSE_SIDE(x_side)  
  
    w ← GET_CHILD(w,r_side)  
    GET_CHILD(w,x_side).color ← BLACK  
    w.color ← RED  
  
    ROTATE(T,w,r_side)  
endif
```

Removing a Key: Code (Cont'd 4)

```
def FIX_REMOVE_RBT_CASE4(T,x):  
    x_side ← CHILDHOOD_SIDE(x)  
    r_side ← REVERSE_SIDE(x_side)  
  
    w ← GET_CHILD(w,r_side)  
    GET_CHILD(w,r_side).color ← BLACK  
    w.color ← x.parent.color  
    x.parent.color ← BLACK  
  
    ROTATE(T,x.parent,x_side)  
endif
```

Removing a Key: Code (Cont'd 5)

```
def FIX_REMOVE_RBT(T, x):  
    while x ≠ T.root and x.color ≠ RED:  
        w ← SIBLING(x)  
        if w = RED:  
            x ← FIX_REMOVE_RBT_CASE1(T, x)  
        else:  
            x_side ← CHILDHOOD_SIDE(x)  
            r_side ← REVERSE_SIDE(x_side)
```

Removing a Key: Code (Cont'd 6)

```
    if GET_CHILD(w, r_size) = RED:
        FIX_REMOVE_RBT_CASE4(T, x)

        return
    else:
        if GET_CHILD(w, x_side) = RED:
            FIX_REMOVE_RBT_CASE3(T, x)
        else:
            FIX_REMOVE_RBT_CASE2(T, x)
        endif
    endif
endwhile
enddef
```